

COMP348 — Document Processing and the Semantic Web

Week 03 Lecture 1: Introduction to Text Classification

Diego Mollá

Department of Computer Science
Macquarie University

COMP348 2018H1

Programme

- 1 What is Text Classification
- 2 Statistical Classification with NLTK and Scikit-Learn
- 3 Advice on Machine Learning

Reading

- NLTK Book Chapter 6 “Learning to Classify Text”

Some Useful Extra Reading

- Ethem Alpaydin (2004). *Introduction to Machine Learning*. MIT Press.

Programme

- 1 What is Text Classification
- 2 Statistical Classification with NLTK and Scikit-Learn
- 3 Advice on Machine Learning

Text Classification

What is Text Classification?

Classify documents into one of a **fixed predetermined** set of categories.

- The number of categories is predetermined.
- The actual categories is predetermined.

Examples

- Spam detection.
- Email filtering.
- Classification of text into genres.
- Classification of names by gender.
- Classification of questions.

Example: Spam Filtering

Distinguish this

Date: Mon, 24 Mar 2008
From: XXX YYY <xxx@yahoo.com>
Subj: Re: Fwd: MSc
To: Mark Dras
<madras@ics.mq.edu.au>

Hi, Thanks for that. It would fit me very well to start 2009, its actually much better for me and I'm planning to finish the project in one year (8 credit points).

from this

Date: Mon, 24 Mar 2008
From: XXX YYY
<xxx@yahoo.co.in>
Subj: HELLO
To: madras@ics.mq.edu.au

HELLO, MY NAME IS STEPHINE IN SEARCH OF A MAN WHO UNDERSTANDS THE MEANING OF LOVE AS TRUST AND FAITH IN EACH OTHER RATHER THAN ONE WHO SEES LOVE AS THE ONLY WAY OF FUN ...

Classification Methods

Manual

- Web portals.
- Wikipedia.

Automatic

Hand coded rules

- e.g. 'Viagra' == SPAM.
- e.g. email filter rules.
- Fragile, breaks on new data.

Supervised learning

- Use an annotated corpus.
- Apply statistical methods.
- Greater flexibility.

Supervised Learning


Given

Training data annotated with class information.

Goal


Build a **model** which will allow classification of new data.

Method




- 1 **Feature extraction:** Convert samples into vectors.
- 2 **Training:** Automatically learn a model.
- 3 **Classification:** Apply the model on new data.

Supervised Learning Example: Bank Customers



(from Alpaydin (2004))


Supervised Learning Example: Family Cars



(from Alpaydin (2004))

Hypothesis Class

- Choice in generality of hypothesis
 - S = Most specific hypothesis: tightest bound
 - G = Most general hypothesis: loosest bound



(from Alpaydin (2004))

Programme


- 1 What is Text Classification
- 2 Statistical Classification with NLTK and Scikit-Learn
 - NLTK
 - Scikit-Learn
- 3 Advice on Machine Learning

Programme

- 1 What is Text Classification
- 2 Statistical Classification with NLTK and Scikit-Learn
 - NLTK
 - Scikit-Learn
- 3 Advice on Machine Learning

NLTK Features

- Statistical classifiers are not able to make sense of text.
- We need to feed them with our interpretation of the text.
- NLTK classifiers expect a dictionary of features and values.



Example of a Simple Feature Extractor

```
def gender_features(word):
    return {'last_letter': word[-1]}
```

Example: Gender Classification

```
>>> import nltk
>>> from nltk.corpus import names
>>> m = names.words('male.txt')
>>> len(m)
2943
>>> f = names.words('female.txt')
>>> len(f)
5001
>>> import random
>>> random.seed(1234) # Fixed random seed to facilitate replicability
>>> names = [(name, 'male') for name in m] +
            [(name, 'female') for name in f]
>>> random.shuffle(names)
>>> def gender_features(word):
    return {'last_letter': word[-1]}
>>> featuresets = [(gender_features(n), g) for n, g in names]
>>> train_set, test_set = featuresets[500:], featuresets[:500]
>>> classifier = nltk.NaiveBayesClassifier.train(train_set)
>>> classifier.classify(gender_features('Neo'))
'male'
>>> classifier.classify(gender_features('Trinity'))
'female'
>>> nltk.classify.accuracy(classifier, test_set)
0.776
>>>
```

Note that the classifier is fed with the gender features and not with the actual names.

The Development Set I

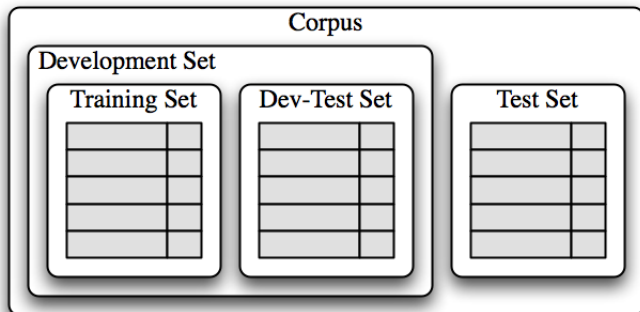
Important

Always test your system with data that has not been used for development (Why ...?)

Development and Test Sets

- Put aside a test set and don't even look at its contents.
- Use the remaining data as a development set.
 - Separate the development set into training and dev-test sets.
 - Use the training set to train the statistical classifiers.
 - Use the dev-test set to fine-tune the classifiers and do error analysis.

The Development Set II



Error Analysis in our Gender Classifier

```
>>> train_names = names[1500:]
>>> devtest_names = names[500:1500]
>>> test_names = names[:500]
>>> train_set = [(gender.features(n), g) for n, g in train_names]
>>> devtest_set = [(gender.features(n), g) for n, g in devtest_names]
>>> test_set = [(gender.features(n), g) for n, g in test_names]
>>> classifier = nltk.NaiveBayesClassifier.train(train_set)
>>> nltk.classify.accuracy(classifier, devtest_set)
0.752
>>> errors = []
>>> for name, tag in devtest_names:
    guess = classifier.classify(gender.features(name))
    if guess != tag:
        errors.append((tag, guess, name))

>>> len(errors)
248
>>> for (tag, guess, name) in sorted(errors):
    print("correct=%0-8s_guess=%0-8s_name=%0-30s" % (tag, guess, name))
```

A Revised Gender Classifier

```
>>> def gender_features2(word):  
    return {'suffix1': word[-1:],  
            'suffix2': word[-2:]}  
  
>>> train_set2 = [(gender_features2(n), g) for n, g in train_names]  
>>> devtest_set2 = [(gender_features2(n), g) for n, g in devtest_names]  
>>> classifier = nltk.NaiveBayesClassifier.train(train_set2)  
>>> nltk.classify.accuracy(classifier, devtest_set2)  
0.77
```

Beware of Over-fitting


If there are many features on a small corpus the system may **over-fit**.

```
>>> def gender_features3(name):
    features = {}
    features['firstletter'] = name[0].lower()
    features['lastletter'] = name[-1]
    for letter in 'abcdefghijklmnopqrstuvwxyz':
        features['count(%s)' % letter] = name.lower().count(letter)
        features['has(%s)' % letter] = (letter in name.lower())
    return features

>>> gender_features3('John')
{'count(u)': 0, 'has(d)': False, 'count(b)': 0, 'count(w)': 0, ...}
>>> train_set3 = [(gender_features3(n), g) for n, g in train_names]
>>> devtest_set3 = [(gender_features3(n), g) for n, g in devtest_names]
>>> classifier = nltk.NaiveBayesClassifier.train(train_set3[:50])
>>> nltk.classify.accuracy(classifier, devtest_set3)
0.698
>>> classifier2b = nltk.NaiveBayesClassifier.train(train_set2[:50])
>>> nltk.classify.accuracy(classifier2b, devtest_set2)
0.7
```

Some types of classifiers are more sensitive to over-fitting than others.

Identifying Over-fitting



Deciding What Features to Keep

NLTK can report on the most informative features.

```
>>> classifier.show_most_informative_features()
```

Most Informative Features

last_letter = 'k'	male : female =	46.2 : 1.0
last_letter = 'a'	female : male =	36.6 : 1.0
last_letter = 'f'	male : female =	15.4 : 1.0
last_letter = 'v'	male : female =	9.9 : 1.0
last_letter = 'p'	male : female =	9.9 : 1.0
last_letter = 'd'	male : female =	9.6 : 1.0
last_letter = 'm'	male : female =	9.1 : 1.0
last_letter = 'o'	male : female =	8.8 : 1.0
last_letter = 'r'	male : female =	6.6 : 1.0
last_letter = 'z'	male : female =	5.1 : 1.0

```
>>>
```

So, if the name ends in 'k', the gender is male 46 times more often than female.

Programme

- 1 What is Text Classification
- 2 Statistical Classification with NLTK and Scikit-Learn
 - NLTK
 - Scikit-Learn
- 3 Advice on Machine Learning

Text Classification in Scikit-Learn

- Scikit-learn includes a large number of statistical classifiers.
- All of the classifiers have a common interface.
- The features of a document set are represented as a matrix.
 - Each row represents a document.
 - Each column represents a feature.
- Scikit-learn provides some useful feature extractors for text:
 - 1 **CountVectorizer** returns a (sparse) matrix of word counts.
 - 2 **TfidfVectorizer** returns a (sparse) matrix of tf.idf values.

Gender Classifier in Scikit-Learn

```
>>> from sklearn.naive_bayes import MultinomialNB
>>> def gender_features(word):
>>>     "Return the ASCII value of the last two characters"
>>>     return [ord(word[-2]), ord(word[-1])]
>>> featuresets = [(gender_features(n), g) for n, g in names]
>>> train_set, test_set = featuresets[500:], featuresets[:500]
>>> train_X, train_y = zip(*train_set)
>>> classifier = MultinomialNB()
>>> classifier.fit(train_X, train_y)
>>> test_X, test_y = zip(*test_set)
>>> classifier.predict(test_X[:5])
array(['female', 'female', 'male', 'female', 'female'],
      dtype='<S6')
```


Programme

- 1 What is Text Classification
- 2 Statistical Classification with NLTK and Scikit-Learn
- 3 Advice on Machine Learning
 - Over-fitting
 - The Dev-test Set

Possible Problems with Machine Learning

- ML methods are typically seen as black boxes.
- Some methods are better than others for specific tasks but people tend to just try several and choose the one with best results.

Possible problems/mistakes you might face

- 1 Train and test are the same dataset (don't do this!).
- 2 The results on the test set are much worse than on the dev-test set.
- 3 The results of both the test set and the training set are bad.
- 4 The train/test partition is not random.
- 5 The results on the test set are good but then the results on your real application problem are bad.

Partition into Training and Testing Set

What's wrong with this partition?

```
from nltk.corpus import names
m = names.words('male.txt')
f = names.words('female.txt')
names = ([ (name, 'm') for name in m]+
          [ (name, 'f') for name in f])
trainset = names[500:]
testset = names[:500]
```

Advice

- 1 Make sure that the train and test sets have no bias.
- 2 Make sure that the train and test sets are representative of your problem.

Partition into Training and Testing Set

What's wrong with this partition?

```
from nltk.corpus import names
m = names.words('male.txt')
f = names.words('female.txt')
names = [(name, 'm') for name in m]+
         [(name, 'f') for name in f]
trainset = names[500:]
testset = names[:500]
```

Advice

- ❶ Make sure that the train and test sets have no bias.
- ❷ Make sure that the train and test sets are representative of your problem.

Randomised Partition


A better partition

```
from nltk.corpus import names
import random
m = names.words('male.txt')
f = names.words('female.txt')
names = [(name, 'm') for name in m]+
        [(name, 'f') for name in f]
random.seed(1234)
random.shuffle(names)
trainset = names[500:]
testset = names[:500]
```

Programme


- 1 What is Text Classification
- 2 Statistical Classification with NLTK and Scikit-Learn
- 3 Advice on Machine Learning
 - Over-fitting
 - The Dev-test Set

Why is machine learning hard?




- There are an **infinite number of curves** that fit the data
 - even more if we don't require the curves to exactly fit (e.g., if we assume there's noise in our data).
- In general, **more data** would help us identify the correct curve better.

Why is machine learning hard?




- There are an **infinite number of curves** that fit the data
 - even more if we don't require the curves to exactly fit (e.g., if we assume there's noise in our data).
- In general, **more data** would help us identify the correct curve better.

Why is machine learning hard?




- There are an **infinite number of curves** that fit the data
 - even more if we don't require the curves to exactly fit (e.g., if we assume there's noise in our data).
- In general, **more data** would help us identify the correct curve better.

Why is machine learning hard?



- There are an **infinite number of curves** that fit the data
 - even more if we don't require the curves to exactly fit (e.g., if we assume there's noise in our data).
- In general, **more data** would help us identify the correct curve better.

Why is machine learning hard?



- There are an **infinite number of curves** that fit the data
 - even more if we don't require the curves to exactly fit (e.g., if we assume there's noise in our data).
- In general, **more data** would help us identify the correct curve better.

The “no free lunch theorem”

- The **“no free lunch theorem”** says there is no single best way to generalise that will be correct in all cases.
 - ⇒ a machine learning algorithm that does well on some problems will do badly on others.
 - ⇒ balancing the trade-off between the **fit to data** and **model complexity** is a central theme in machine learning.
- Even so, in practice there are machine learning algorithms that do well on broad classes of problems.
- But it's important to understand the problem you are trying to solve as well as possible.


The “no free lunch theorem”

- The **“no free lunch theorem”** says there is no single best way to generalise that will be correct in all cases.
 - ⇒ a machine learning algorithm that does well on some problems will do badly on others.
 - ⇒ balancing the trade-off between the **fit to data** and **model complexity** is a central theme in machine learning.
- Even so, in practice there are machine learning algorithms that do well on broad classes of problems.
- But it's important to understand the problem you are trying to solve as well as possible.

The “no free lunch theorem”

- The **“no free lunch theorem”** says there is no single best way to generalise that will be correct in all cases.
 - ⇒ a machine learning algorithm that does well on some problems will do badly on others.
 - ⇒ balancing the trade-off between the **fit to data** and **model complexity** is a central theme in machine learning.
- Even so, in practice there are machine learning algorithms that do well on broad classes of problems.
- But it's important to understand the problem you are trying to solve as well as possible.

Over-fitting the Training Data



- Over-fitting occurs when an algorithm learns a function that is fitting noise in the data.
- Diagnostic of over-fitting: performance on training data is much higher than performance on dev or test data.


Programme

- 1 What is Text Classification
- 2 Statistical Classification with NLTK and Scikit-Learn
- 3 Advice on Machine Learning
 - Over-fitting
 - The Dev-test Set

Fine-tuning the System


- Many ML algorithms have parameters that adjust the tendency to learn from the data.
- A reasonable approach to find the best parameters is to try many options and select the best one.
- For example, K-NN uses the number of neighbours (K) as a parameter.
 - Smaller K learns from training data more effectively but it may learn the noise in the training data (over-fitting).

Example of the Impact of K with the Names Corpus



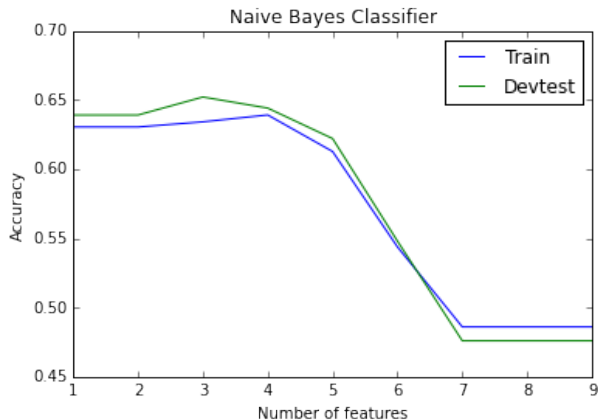
See related notebook for the code.

Example of the Impact of Features with the Names Corpus



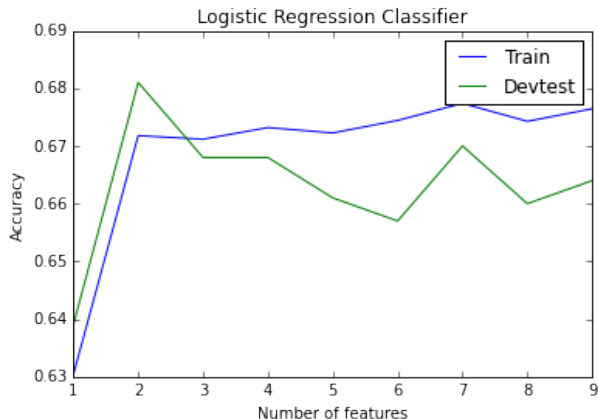
See related notebook for the code.

Example of the Impact of Features with the Names Corpus



See related notebook for the code.

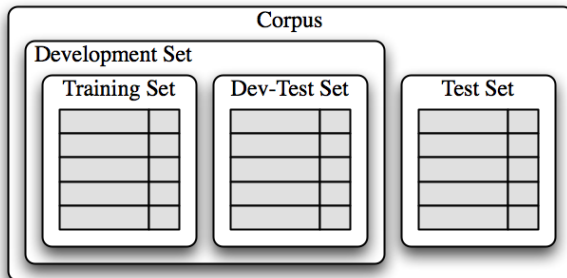
Example of the Impact of Features with the Names Corpus



See related notebook for the code.

Introducing the Dev-test Set

- You need to use data that is not in the training set nor in the test set to fine-tune the system.
- It is therefore advisable to have **three partitions**.
- The test set is used only for the final test, not for fine-tuning parameters.



Take-home Messages

- 1 Explain and demonstrate the need for separate training, dev-test, and test set.
- 2 Implement feature extractors for statistical classifiers in NLTK and Scikit-Learn.
- 3 Use NLTK's and Scikit-Learn's statistical classifiers.
- 4 Detect over-fitting.

What's Next

Week 4

- Text Classification II.

Reading

- NLTK Chapter 6.